

```
In [130]: ## Importing Libraries

import numpy as np # Linear algebra operations
import pandas as pd # Data processing and analysis
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf

from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, confusion_matrix, roc_auc
from sklearn.neighbors import KNeighborsClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn import tree
from sklearn.svm import SVC
from tensorflow import keras
from tensorflow.keras import layers, Sequential
from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
from tensorflow.keras.layers import Dense, Dropout, BatchNormalization, LSTM
from sklearn.ensemble import GradientBoostingClassifier
from sklearn.naive_bayes import GaussianNB

import warnings
warnings.filterwarnings("ignore")
```

```
In [3]: ## Upload dataset

df = pd.read_csv('/Users/serenaygoler/heart disease.csv')

df.head() # Displays the first 5 rows.
```

```
Out[3]:
```

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxHR
0	40	M	ATA	140	289	0	Normal	172
1	49	F	NAP	160	180	0	Normal	156
2	37	M	ATA	130	283	0	ST	98
3	48	F	ASY	138	214	0	Normal	108
4	54	M	NAP	150	195	0	Normal	122

```
In [6]: df.tail() # Display the last 5 rows.
```

Out [6]:

	Age	Sex	ChestPainType	RestingBP	Cholesterol	FastingBS	RestingECG	MaxH
913	45	M	TA	110	264	0	Normal	13
914	68	M	ASY	144	193	1	Normal	14
915	57	M	ASY	130	131	0	Normal	17
916	57	F	ATA	130	236	0	LVH	17
917	38	M	NAP	138	175	0	Normal	17

In [8]: `df.info()` # Prints name and type of variables, number of observations, and c

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 918 entries, 0 to 917
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Age                    918 non-null   int64
1   Sex                    918 non-null   object
2   ChestPainType          918 non-null   object
3   RestingBP              918 non-null   int64
4   Cholesterol            918 non-null   int64
5   FastingBS              918 non-null   int64
6   RestingECG             918 non-null   object
7   MaxHR                  918 non-null   int64
8   ExerciseAngina         918 non-null   object
9   Oldpeak                918 non-null   float64
10  ST_Slope               918 non-null   object
11  HeartDisease           918 non-null   int64
dtypes: float64(1), int64(6), object(5)
memory usage: 86.2+ KB
```

In [10]: `df.shape` # Displays the number of rows and columns in the dataset.

Out[10]: (918, 12)

In [12]: `df.isna().sum()` # Counts missing values in each column.

```
Out[12]: Age                0
Sex                  0
ChestPainType        0
RestingBP            0
Cholesterol          0
FastingBS            0
RestingECG           0
MaxHR                0
ExerciseAngina        0
Oldpeak              0
ST_Slope             0
HeartDisease         0
dtype: int64
```

In [14]: `df.duplicated().sum()` # Counts the number of duplicate rows.

Out [14]: 0

In [16]: `## Provides summary statistics for numeric columns, rounded to 2 decimals and
df.describe().round(2).T`

Out [16]:

	count	mean	std	min	25%	50%	75%	max
Age	918.0	53.51	9.43	28.0	47.00	54.0	60.0	77.0
RestingBP	918.0	132.40	18.51	0.0	120.00	130.0	140.0	200.0
Cholesterol	918.0	198.80	109.38	0.0	173.25	223.0	267.0	603.0
FastingBS	918.0	0.23	0.42	0.0	0.00	0.0	0.0	1.0
MaxHR	918.0	136.81	25.46	60.0	120.00	138.0	156.0	202.0
Oldpeak	918.0	0.89	1.07	-2.6	0.00	0.6	1.5	6.2
HeartDisease	918.0	0.55	0.50	0.0	0.00	1.0	1.0	1.0

In [18]: `# Count how many Cholesterol values are zero
chol_zero_count = (df["Cholesterol"] == 0).sum()

Count how many RestingBP values are zero
bp_zero_count = (df["RestingBP"] == 0).sum()

print(f"Number of Cholesterol values equal to 0: {chol_zero_count}")
print(f"Number of RestingBP values equal to 0: {bp_zero_count}")`

Number of Cholesterol values equal to 0: 172

Number of RestingBP values equal to 0: 1

In [20]: `# Cross-tabulate Cholesterol = 0 with HeartDisease status
import pandas as pd

zero_chol = df[df["Cholesterol"] == 0]
ct = pd.crosstab(zero_chol["HeartDisease"], zero_chol["Cholesterol"])
print(ct)`

Cholesterol	0
HeartDisease	
0	20
1	152

In [22]: `## Filters out rows where Cholesterol equals zero and returns summary statistics
print(df[df["Cholesterol"] != 0]["Cholesterol"].describe())`

```

count      746.000000
mean       244.635389
std        59.153524
min        85.000000
25%        207.250000
50%        237.000000
75%        275.000000
max        603.000000
Name: Cholesterol, dtype: float64

```

```

In [24]: # With zeros included
print("=== With Zero values Included ===")
print(df.groupby("HeartDisease")["Cholesterol"].describe())

# Zeros removed
print("\n=== With zero values removed ===")
print(df[df["Cholesterol"] != 0].groupby("HeartDisease")["Cholesterol"].desc

```

```

=== With Zero values Included ===

```

	count	mean	std	min	25%	50%	75%	max
HeartDisease								
0	410.0	227.121951	74.634659	0.0	197.25	227.0	266.75	564.0
1	508.0	175.940945	126.391398	0.0	0.00	217.0	267.00	603.0

```

=== With zero values removed ===


```

	count	mean	std	min	25%	50%	75%	max
HeartDisease								
0	390.0	238.769231	55.394617	85.0	203.0	231.5	269.00	564.0
1	356.0	251.061798	62.462713	100.0	212.0	246.0	283.25	603.0

```

In [26]: # Plot the cholesterol distribution for HeartDisease=0 and HeartDisease=1 us
plt.figure(figsize=(10,6))
sns.kdeplot(df[df["HeartDisease"]==0]["Cholesterol"], label="HeartDisease=0")
sns.kdeplot(df[df["HeartDisease"]==1]["Cholesterol"], label="HeartDisease=1")
plt.legend()
plt.title("Cholesterol Density Distribution")
plt.show()

```



```
In [28]: # This block cleans the dataset by:
# 1. Removing rows where RestingBP = 0 (unrealistic values).
# 2. Calculating group-wise medians of Cholesterol (by HeartDisease) excluding
# 3. Replacing Cholesterol values of zero with the corresponding group median
# 4. Checking that no zero values remain.
# 5. Displaying summary statistics of Cholesterol by HeartDisease after cleaning

df_clean = df.copy()
df_clean = df_clean[df_clean["RestingBP"] != 0].copy()

medians = (
    df_clean[df_clean["Cholesterol"] != 0]
    .groupby("HeartDisease")["Cholesterol"]
    .median()
)


mask_zero = df_clean["Cholesterol"] == 0
df_clean["Cholesterol"] = df_clean["Cholesterol"].astype(float)
df_clean.loc[mask_zero, "Cholesterol"] = (
    df_clean.loc[mask_zero, "HeartDisease"].map(medians)
)

print("Remaining zeros:", (df_clean["Cholesterol"] == 0).sum())
print(df_clean.groupby("HeartDisease")["Cholesterol"].describe())
```

Remaining zeros: 0

	count	mean	std	min	25%	50%	75%	m
ax								
HeartDisease								
0	410.0	238.414634	54.045994	85.0	204.0	231.5	266.75	56
4.0								
1	507.0	249.554241	52.370323	100.0	225.0	246.0	267.00	60
3.0								

```
In [30]: # KDE plot – distribution comparison after median imputation
plt.figure(figsize=(10,6))
sns.kdeplot(df_clean[df_clean["HeartDisease"]==0]["Cholesterol"], label="HeartDisease=0")
sns.kdeplot(df_clean[df_clean["HeartDisease"]==1]["Cholesterol"], label="HeartDisease=1")
plt.title("Cholesterol Distribution (After Median Imputation)")
plt.xlabel("Cholesterol")
plt.ylabel("Density")
plt.legend()
plt.show()
```



```
In [32]: ## Provides summary statistics for numeric columns for clean data, rounded to 2 decimal places
df_clean.describe().round(2).T
```

Out [32]:

	count	mean	std	min	25%	50%	75%	max
Age	917.0	53.51	9.44	28.0	47.0	54.0	60.0	77.0
RestingBP	917.0	132.54	18.00	80.0	120.0	130.0	140.0	200.0
Cholesterol	917.0	244.57	53.39	85.0	214.0	246.0	267.0	603.0
FastingBS	917.0	0.23	0.42	0.0	0.0	0.0	0.0	1.0
MaxHR	917.0	136.79	25.47	60.0	120.0	138.0	156.0	202.0
Oldpeak	917.0	0.89	1.07	-2.6	0.0	0.6	1.5	6.2
HeartDisease	917.0	0.55	0.50	0.0	0.0	1.0	1.0	1.0


In [34]: *# Count negative Oldpeak values*

```
neg_oldpeak_count = (df["Oldpeak"] < 0).sum()
print(f"Number of negative Oldpeak values: {neg_oldpeak_count}")
```

Number of negative Oldpeak values: 13

In [36]: *# Stripplot showing distribution of Oldpeak values by HeartDisease, with ref*

```
plt.figure(figsize=(8,5))
sns.stripplot(x="HeartDisease", y="Oldpeak", data=df, jitter=True, palette="
plt.axhline(0, color="red", linestyle="--")
plt.title("Oldpeak Values by HeartDisease (individual points)")
plt.show()
```



```

In [38]: # Plot numeric feature distributions by target, two-at-a-time
num_cols = df_clean.select_dtypes(include="number").columns.drop("HeartDisease")
cols = list(num_cols)

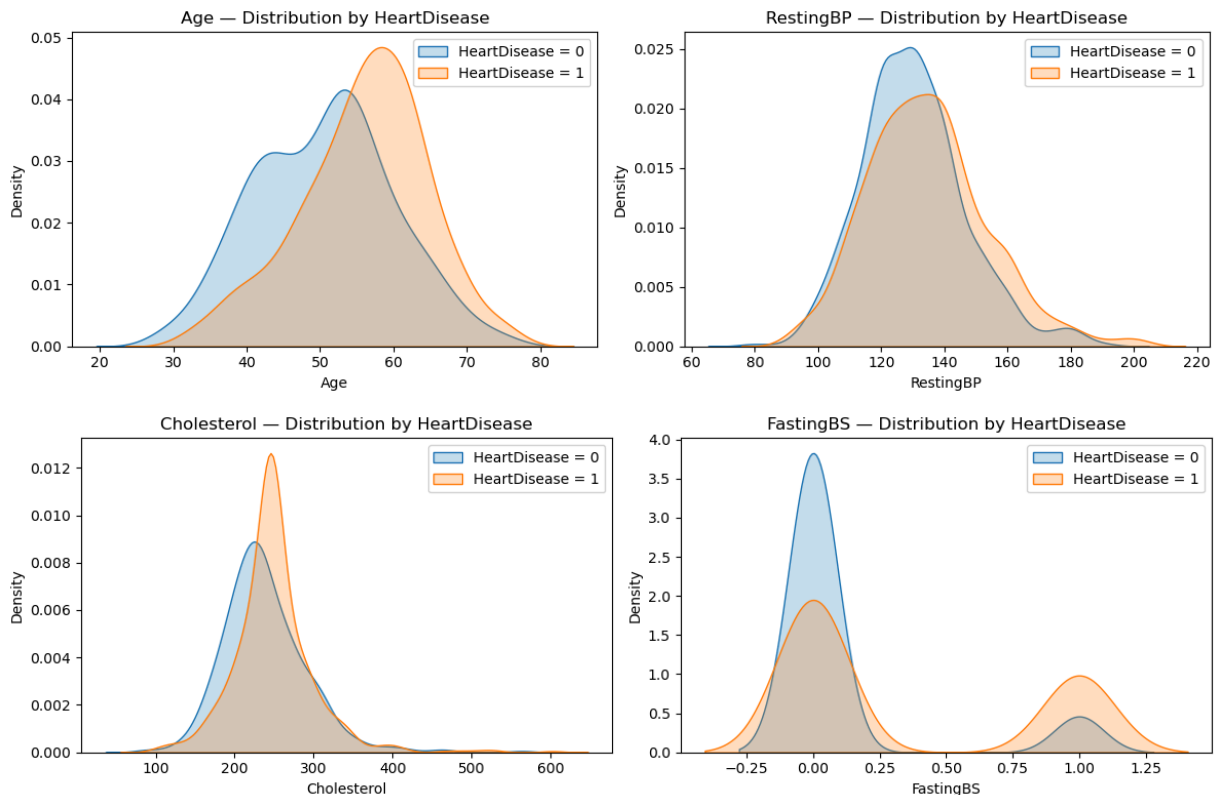
for i in range(0, len(cols), 2):
    pair = cols[i:i+2] # up to 2 columns per figure


    fig, axes = plt.subplots(1, len(pair), figsize=(12, 4))
    if len(pair) == 1:
        axes = [axes] # make iterable if only one axis

    for ax, col in zip(axes, pair):
        sns.kdeplot(
            df_clean.loc[df_clean["HeartDisease"] == 0, col].dropna(),
            label="HeartDisease = 0", fill=True, ax=ax
        )
        sns.kdeplot(
            df_clean.loc[df_clean["HeartDisease"] == 1, col].dropna(),
            label="HeartDisease = 1", fill=True, ax=ax
        )
        ax.set_title(f"{col} — Distribution by HeartDisease")
        ax.set_xlabel(col); ax.set_ylabel("Density")
        ax.legend()

plt.tight_layout()
plt.show()


```





```
In [40]: # Select only numerical columns and to check correlation
num_cols = df_clean.select_dtypes(include=[np.number]).columns

plt.figure(figsize=(8,6))
sns.heatmap(df_clean[num_cols].corr(), annot=True, cmap="viridis", fmt=".2f")
plt.title("Correlation Heatmap (numeric features)")
plt.show()
```



```
In [42]: # Distribution of categorical variables by the target variable
cat_cols = ["Sex", "ChestPainType", "FastingBS", "RestingECG", "ExerciseAngi"]
fig, axes = plt.subplots(nrows=3, ncols=2, figsize=(16, 14))
axes = axes.flatten()
```

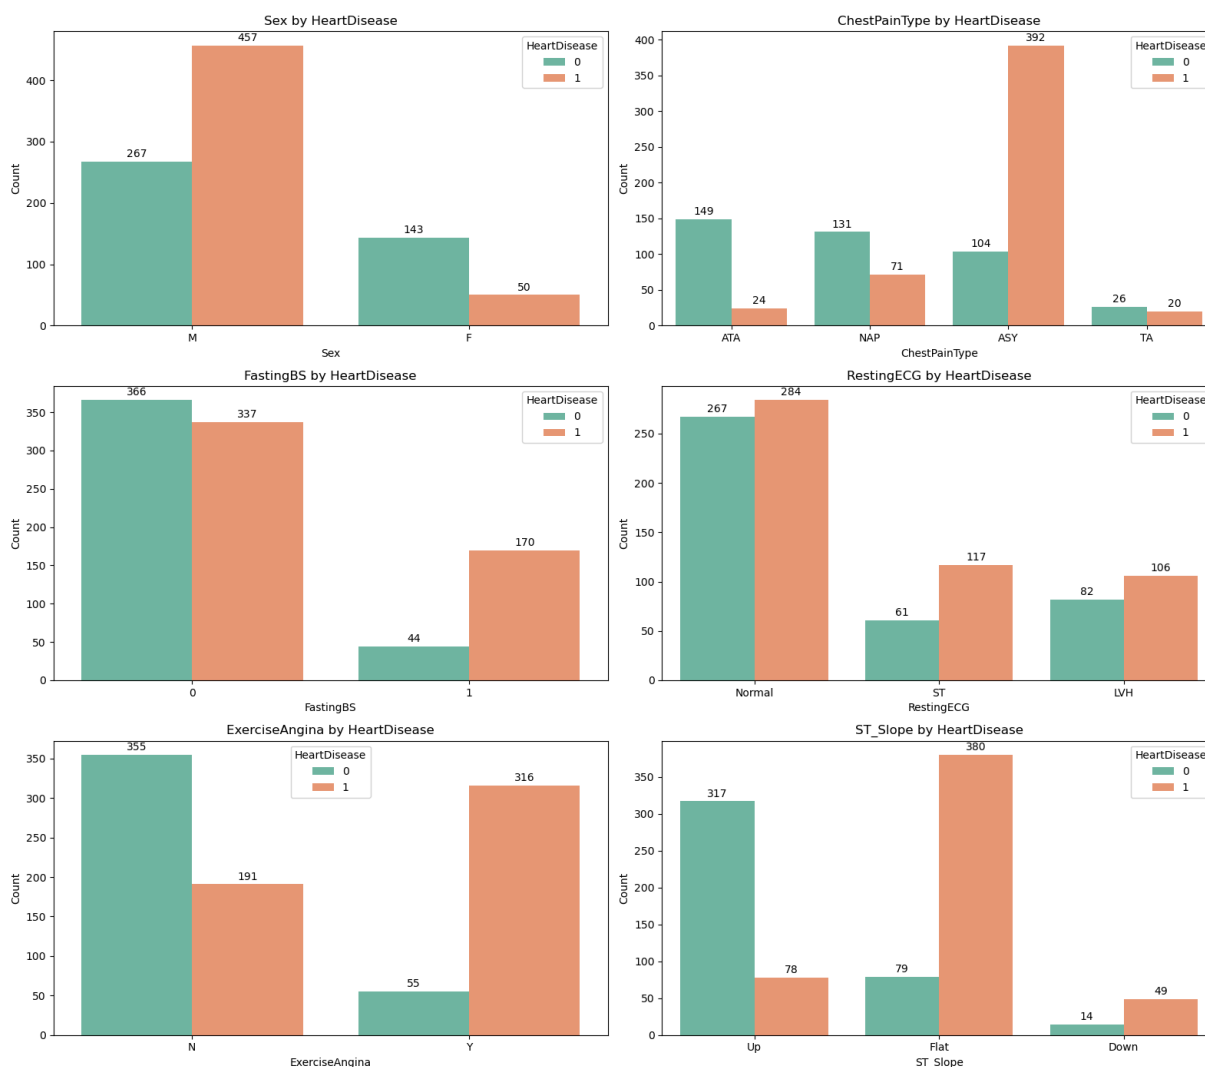
```

for ax, col in zip(axes, cat_cols):
    g = sns.countplot(data=df_clean, x=col, hue="HeartDisease", palette="Set2")
    ax.set_title(f"{col} by HeartDisease")
    ax.set_xlabel(col); ax.set_ylabel("Count")
    # label name
    for c in g.containers:
        g.bar_label(c, padding=2, fmt="%.0f")

# Remove extra axes
for ax in axes[len(cat_cols):]:
    fig.delaxes(ax)

plt.tight_layout()
plt.show()

```



```

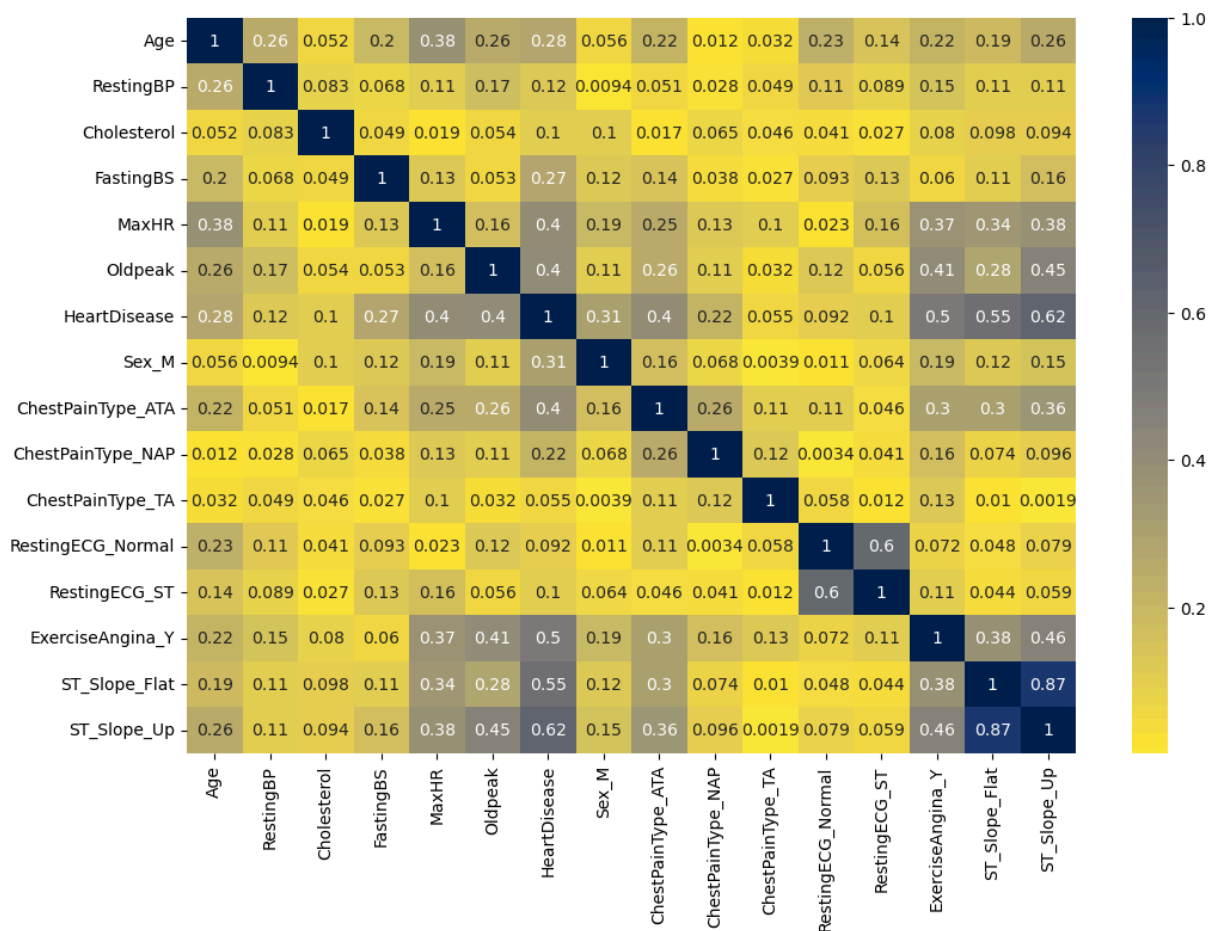
In [44]: # One-hot encoding was applied to transform categorical variables into dummy
DUMMY = pd.get_dummies(df_clean, drop_first=True)
DUMMY.head()

```

```
Out [44]:
```

	Age	RestingBP	Cholesterol	FastingBS	MaxHR	Oldpeak	HeartDisease	Sex_M
0	40	140	289.0	0	172	0.0	0	True
1	49	160	180.0	0	156	1.0	1	False
2	37	130	283.0	0	98	0.0	0	True
3	48	138	214.0	0	108	1.5	1	False
4	54	150	195.0	0	122	0.0	0	True

```
In [46]: ## Compute absolute pairwise correlations (after one-hot encoding) and visualize
correlations = abs(DUMMY.corr())
plt.figure(figsize=(12,8))
sns.heatmap(correlations, annot=True, cmap="cividis_r")
plt.show()
```



```
In [48]: # Create a copy of the cleaned dataset
codedf = df_clean.copy()

# 1) Convert binary categorical columns into 0/1 format
if codedf['Sex'].dtype == 'object':
    codedf['Sex'] = codedf['Sex'].str.strip().map({'F': 0, 'M': 1}).astype('int')

if codedf['ExerciseAngina'].dtype == 'object':
    codedf['ExerciseAngina'] = codedf['ExerciseAngina'].str.strip().map({'N': 0, 'Y': 1}).astype('int')
```

```
# (If they are already boolean True/False, convert them to integers)
for col in ['Sex', 'ExerciseAngina']:
    if codedf[col].dtype == 'bool':
        codedf[col] = codedf[col].astype(int)

# 2) Apply one-hot encoding for multi-class categorical columns
to_onehot = ['ChestPainType', 'RestingECG', 'ST_Slope']
codedf = pd.get_dummies(codedf, columns=to_onehot, drop_first=True)

# Convert any remaining boolean columns into 0/1 integers
for col in codedf.select_dtypes(include='bool').columns:
    codedf[col] = codedf[col].astype(int)

codedf.dtypes
```

```
Out[48]: Age                int64
Sex                Int64
RestingBP          int64
Cholesterol        float64
FastingBS          int64
MaxHR              int64
ExerciseAngina     Int64
Oldpeak            float64
HeartDisease       int64
ChestPainType_ATA  int64
ChestPainType_NAP  int64
ChestPainType_TA   int64
RestingECG_Normal  int64
RestingECG_ST      int64
ST_Slope_Flat      int64
ST_Slope_Up        int64
dtype: object
```

```
In [50]: # Standardize continuous variables (mean = 0, std = 1)
# This ensures that all numeric predictors are on the same scale,
# which is especially important for distance-based algorithms (e.g., KNN, SVM)
numcolsc = ['Age', 'RestingBP', 'Cholesterol', 'MaxHR', 'Oldpeak']
scaler = StandardScaler()
codedf[numcolsc] = scaler.fit_transform(codedf[numcolsc])

codedf.head()
```

```
Out[50]:
```

	Age	Sex	RestingBP	Cholesterol	FastingBS	MaxHR	ExerciseAngina	Oldpeak
0	-1.432206	1	0.414627	0.832639	0	1.383339	0	-0.8
1	-0.478057	0	1.526360	-1.210238	0	0.754736	0	0.7
2	-1.750256	1	-0.141240	0.720187	0	-1.523953	0	-0.8
3	-0.584074	0	0.303453	-0.573010	0	-1.131075	1	0.5
4	0.052026	1	0.970493	-0.929108	0	-0.581047	0	-0.8

Machine Learning

```
In [53]: # Split the dataset into features (X) and target (y)
X = codedf.drop(columns=["HeartDisease"])
y = codedf["HeartDisease"]

# Train-test split: 80% training, 20% testing
# Stratify ensures the target class distribution (0/1) is preserved in both
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

```
In [55]: X_test.shape, y_test.shape # Check the shape of the test sets
```

```
Out[55]: ((184, 15), (184,))
```

Logistic Regression

```
In [58]: # Logistic Regression Model
# max_iter=1000 ensures convergence during optimization
logreg = LogisticRegression(max_iter=1000)

# Train the model on the training set
logreg.fit(X_train, y_train)

# Make predictions on the test set
y_pred = logreg.predict(X_test)

# Calculate accuracy of the model
logregAcc = accuracy_score(y_test, y_pred)
logregAcc
```

```
Out[58]: 0.8858695652173914
```

```
In [60]: # Generate a detailed classification report
# Includes precision, recall, f1-score, and support for each class
print("\nClassification Report:\n", classification_report(y_test, y_pred))

# Predict probabilities for the positive class (1 = Heart Disease)
y_proba = logreg.predict_proba(X_test)[:, 1]

# Calculate the ROC AUC score to evaluate the model's discriminative ability
print("ROC AUC:", roc_auc_score(y_test, y_proba))
```

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.87	0.87	82
1	0.89	0.90	0.90	102
accuracy			0.89	184
macro avg	0.88	0.88	0.88	184
weighted avg	0.89	0.89	0.89	184


ROC AUC: 0.9423720707795313

```
In [62]: from sklearn.metrics import roc_curve, roc_auc_score

# Probability predictions for positive class
y_proba = logreg.predict_proba(X_test)[:,-1]


# ROC curve values
fpr, tpr, thresholds = roc_curve(y_test, y_proba)
roc_auc = roc_auc_score(y_test, y_proba)

# ROC curve plot
plt.figure(figsize=(6, 6))
plt.plot(fpr, tpr, color='blue', lw=2,
         label='Logistic Regression (AUC = %0.2f)' % roc_auc)
plt.plot([0, 1], [0, 1], color='darkblue', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Logistic Regression')
plt.legend(loc="lower right")
plt.show()
```



```
In [64]: ## Confusion Matrix
cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", cbar=False,
            xticklabels=['No Heart Disease', 'Heart Disease'],
            yticklabels=['No Heart Disease', 'Heart Disease'])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix - Logistic Regression")
plt.show()
```



```
In [66]: # Create a dataframe of Logistic Regression coefficients
# This shows the direction (+/-) and relative magnitude of each feature's effect
# Positive coefficients → increase likelihood of heart disease
# Negative coefficients → decrease likelihood of heart disease

coefficients = pd.DataFrame({
    'Feature': X_train.columns,
    'Coefficient': logreg.coef_[0]
}).sort_values(by='Coefficient', ascending=False)

coefficients
```


Out [66]:

	Feature	Coefficient
1	Sex	1.358389
4	FastingBS	1.156245
13	ST_Slope_Flat	0.948000
6	ExerciseAngina	0.828883
7	Oldpeak	0.322847
3	Cholesterol	0.074785
0	Age	0.039656
2	RestingBP	-0.008738
12	RestingECG_ST	-0.173987
5	MaxHR	-0.268010
11	RestingECG_Normal	-0.309394
14	ST_Slope_Up	-1.280401
10	ChestPainType_TA	-1.329603
9	ChestPainType_NAP	-1.485096
8	ChestPainType_ATA	-1.554887

SUPPORT VECTOR MACHINE

```
In [69]: # Linear SVM
svm_linear = SVC(kernel="linear", probability=True, random_state=1)
svm_linear.fit(X_train, y_train)
svm_linearAcc = accuracy_score(y_test, svm_linear.predict(X_test))
print("Linear SVM Accuracy:", svm_linearAcc)

# RBF SVM (non-linear)
svm_rbf = SVC(kernel="rbf", probability=True, random_state=1)
svm_rbf.fit(X_train, y_train)
svm_rbfAcc = accuracy_score(y_test, svm_rbf.predict(X_test))
print("RBF SVM Accuracy:", svm_rbfAcc)
```

Linear SVM Accuracy: 0.8586956521739131

RBF SVM Accuracy: 0.8858695652173914

```
In [71]: print("\nClassification Report:\n", classification_report(y_test, y_pred))

# Note:
# The ROC curve and confusion matrix visuals are not repeated here for SVM,
# as their performance and outputs were nearly identical to Logistic Regress
```

```

Classification Report:
              precision    recall  f1-score   support

     0       0.88        0.87        0.87         82
     1       0.89        0.90        0.90        102

 accuracy          0.89         184
 macro avg         0.88         184
 weighted avg      0.89         184

```

DECISION TREE

```

In [74]: # Build and train a Decision Tree model
clf = tree.DecisionTreeClassifier(random_state=0) # reproducibility ensured
clf.fit(X_train, y_train) # fit the model to training data

# Make predictions on the test set
y_pred = clf.predict(X_test)

# Calculate accuracy on the test set
clfAcc = accuracy_score(y_test, y_pred)
clfAcc

```

Out[74]: 0.7663043478260869

```

In [76]: # Print classification metrics and ROC AUC for the Decision Tree
print("\nClassification Report:\n", classification_report(y_test, y_pred))

# Predict probabilities for the positive class (heart disease = 1)
y_proba = clf.predict_proba(X_test)[:, 1]

# Calculate and print ROC AUC score
print("ROC AUC:", roc_auc_score(y_test, y_proba))

```

```

Classification Report:
              precision    recall  f1-score   support

     0       0.73        0.74        0.74         82
     1       0.79        0.78        0.79        102

 accuracy          0.77         184
 macro avg         0.76         184
 weighted avg      0.77         184

```

ROC AUC: 0.7641080822572931

```

In [78]: # Predict probability estimates for the positive class (heart disease = 1)
y_proba_tree = clf.predict_proba(X_test)[:, 1]

# Compute ROC curve values
fpr, tpr, thresholds = roc_curve(y_test, y_proba_tree)

# Calculate AUC (Area Under the Curve)
roc_auc_tree = roc_auc_score(y_test, y_proba_tree)


```

```

# Plot ROC curve for Decision Tree
plt.figure(figsize=(6, 6))
plt.plot(fpr, tpr, color='darkgreen', lw=2,
         label='Decision Tree (AUC = %0.2f)' % roc_auc_tree)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Decision Tree')
plt.legend(loc="lower right")
plt.show()

# Print final AUC value
print("ROC AUC (Decision Tree):", roc_auc_tree)

```



ROC AUC (Decision Tree): 0.7641080822572931

```

In [80]: # Confusion Matrix
cm = confusion_matrix(y_test, y_pred)


plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Greens", cbar=False,

```

```

xticklabels=['No Heart Disease','Heart Disease'],
yticklabels=['No Heart Disease','Heart Disease'])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix - Decision Tree")
plt.show()

```



RANDOM FOREST

```

In [83]: # 1) Build and train the model
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# 2) Predictions on the test set
y_pred = rf.predict(X_test)
y_proba = rf.predict_proba(X_test)[:, 1] # probability for the positive class

# 3) Evaluation
rfAcc = accuracy_score(y_test, y_pred) # store accuracy
roc_auc_rf = roc_auc_score(y_test, y_proba) # store AUC
report = classification_report(y_test, y_pred) # precision/recall/F1 per class

print("Random Forest Accuracy:", rfAcc)
print("\nClassification Report:\n", report)
print("ROC AUC:", roc_auc_rf)

```

Random Forest Accuracy: 0.8804347826086957

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.85	0.86	82
1	0.88	0.90	0.89	102
accuracy			0.88	184
macro avg	0.88	0.88	0.88	184
weighted avg	0.88	0.88	0.88	184

ROC AUC: 0.93824725011956

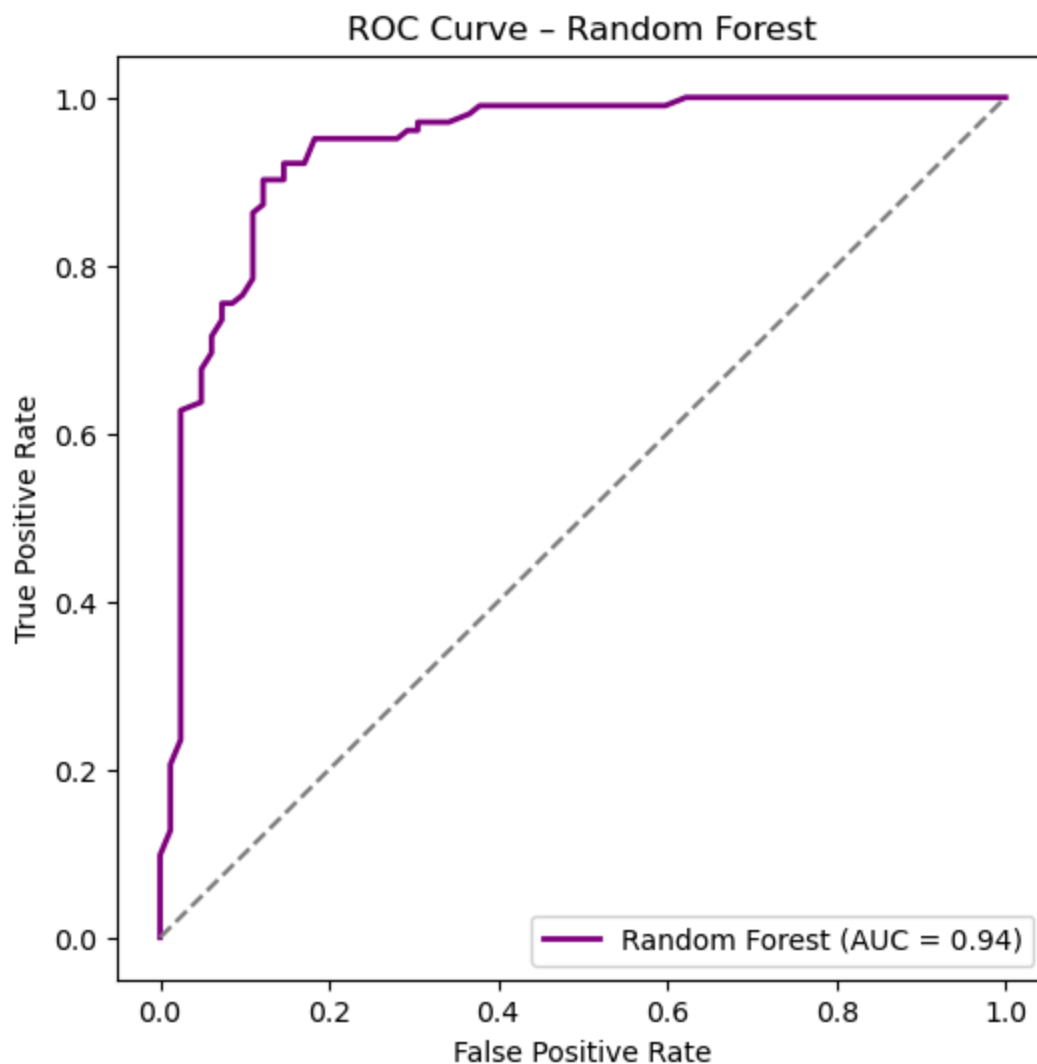
In [85]: *# ROC Curve for Random Forest*

```

y_proba_rf = rf.predict_proba(X_test)[: , 1]
fpr, tpr, _ = roc_curve(y_test, y_proba_rf)
auc_rf = roc_auc_score(y_test, y_proba_rf)

plt.figure(figsize=(6,6))
plt.plot(fpr, tpr, color="purple", lw=2, label=f"Random Forest (AUC = {auc_r
plt.plot([0,1],[0,1], '--', color='gray')
plt.xlabel("False Positive Rate"); plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Random Forest"); plt.legend(loc="lower right")
plt.show()


```



```
In [87]: # Confusion Matrix
cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Purples", cbar=False,
            xticklabels=["Predicted 0 (No Disease)", "Predicted 1 (Disease)",
                        "Actual 0 (No Disease)", "Actual 1 (Disease)"])


plt.title("Confusion Matrix - Random Forest")
plt.ylabel("Actual")
plt.xlabel("Predicted")
plt.tight_layout()
plt.show()
```



```
In [89]: # Extract and sort feature importances from the trained Random Forest model
fi = pd.Series(rf.feature_importances_, index=X_train.columns).sort_values()

# Select the top 10 most important features
top = fi.tail(10)

# Plot horizontal bar chart of feature importances
plt.figure(figsize=(7,5))
top.plot(kind="barh")
plt.title("Top Feature Importances - Random Forest")
plt.xlabel("Importance")
plt.tight_layout()
plt.show()
```



KNN

```
In [92]: # Define the range of k values to test
k_values = range(1, 21)
cv_scores = []

# Loop through each k and perform 5-fold cross-validation
for k in k_values:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X_train, y_train, cv=5, scoring='accuracy')
    cv_scores.append(scores.mean())

# Identify the k with the highest mean accuracy
best_k = k_values[cv_scores.index(max(cv_scores))]

print("Best k:", best_k)
```

Best k: 10

```
In [94]: # Building a model using KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors = 10)
knn.fit(X_train, y_train)

y_pred = knn.predict(X_test)
knnAcc = accuracy_score(y_test, y_pred)
knnAcc
```

Out[94]: 0.8695652173913043


```
In [96]: # Generate probability predictions for the positive class
y_proba = knn.predict_proba(X_test)[: , 1]

# Evaluate the model with ROC AUC and classification report
print("ROC AUC:", roc_auc_score(y_test, y_proba))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

ROC AUC: 0.9324485891917742


Classification Report:

	precision	recall	f1-score	support
0	0.83	0.89	0.86	82
1	0.91	0.85	0.88	102
accuracy			0.87	184
macro avg	0.87	0.87	0.87	184
weighted avg	0.87	0.87	0.87	184

```
In [98]: # Probability predictions for the positive class
y_proba_knn = knn.predict_proba(X_test)[: , 1]


# ROC curve and AUC
fpr, tpr, _ = roc_curve(y_test, y_proba_knn)
auc_knn = roc_auc_score(y_test, y_proba_knn)

# Plot ROC curve
plt.figure(figsize=(6,6))
plt.plot(fpr, tpr, color='orange', lw=2, label="KNN (AUC = %0.2f)" % auc_knn)
plt.plot([0,1],[0,1], '--', color='gray')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - KNN (k=10)")
plt.legend(loc="lower right")
plt.show()
```



```
In [100... # Confusion Matrix
cm = confusion_matrix(y_test, y_pred)

plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt="d", cmap="Oranges", cbar=False,
            xticklabels=['No Heart Disease', 'Heart Disease'],
            yticklabels=['No Heart Disease', 'Heart Disease'])
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.title("Confusion Matrix - KNN")
plt.show()
```



GRADIENT BOOSTING CLASSIFIER

```
In [103... # 1) Train
gboost = GradientBoostingClassifier(
    random_state=42,      # reproducibility
    learning_rate=0.05,   # mild shrinkage
    n_estimators=300,     # a bit larger to stabilize
    max_depth=2           # shallow trees (prevents overfitting on small data)
)
gboost.fit(X_train, y_train)

# 2) Predict class labels and probabilities
y_pred_gb = gboost.predict(X_test)
y_proba_gb = gboost.predict_proba(X_test)[:, 1]

# 3) Metrics (keep them all!)
gboostAcc = accuracy_score(y_test, y_pred_gb)
gboostAUC = roc_auc_score(y_test, y_proba_gb)

print("Gradient Boosting Accuracy:", gboostAcc)
print("Gradient Boosting ROC AUC:", gboostAUC)
print("\nClassification Report (GBoost):\n", classification_report(y_test, y_pred_gb))
print("Confusion Matrix (GBoost):\n", confusion_matrix(y_test, y_pred_gb))

# 4) ROC curve
fpr, tpr, _ = roc_curve(y_test, y_proba_gb)
plt.figure(figsize=(6,6))
plt.plot(fpr, tpr, lw=2, color="deeppink",
        label=f"GBoost (AUC = {gboostAUC:.2f})")
plt.plot([0,1], [0,1], '--', color='lightgray')
plt.xlabel("False Positive Rate")
```

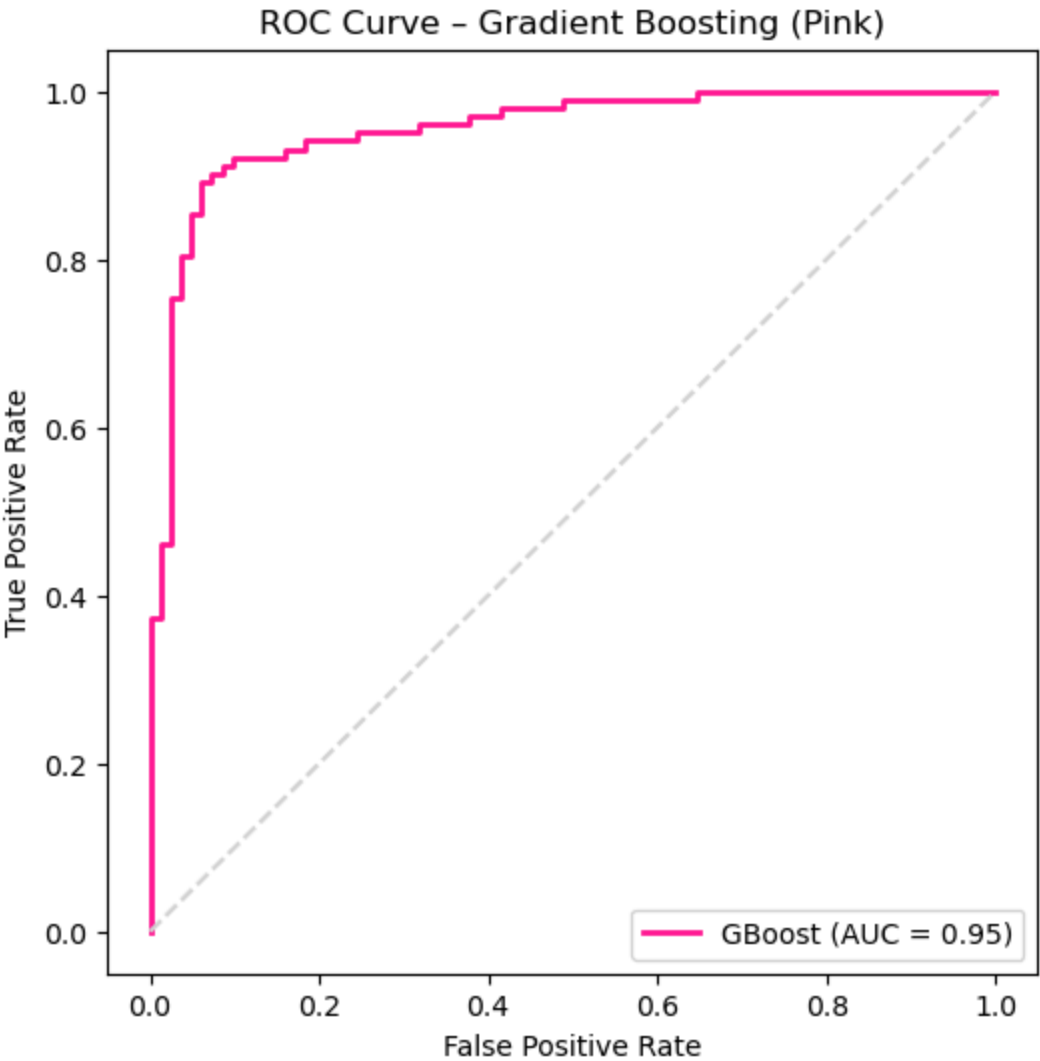
```
plt.ylabel("True Positive Rate")
plt.title("ROC Curve - Gradient Boosting (Pink)")
plt.legend(loc="lower right")
plt.show()
```

Gradient Boosting Accuracy: 0.907608695652174
Gradient Boosting ROC AUC: 0.9549258727881397

Classification Report (GBoost):


	precision	recall	f1-score	support
0	0.88	0.91	0.90	82
1	0.93	0.90	0.92	102
accuracy			0.91	184
macro avg	0.91	0.91	0.91	184
weighted avg	0.91	0.91	0.91	184

Confusion Matrix (GBoost):
[[75 7]
[10 92]]



```
In [105... # --- Confusion Matrix Visualization for GBoost (Pink tones) ---
cm = confusion_matrix(y_test, y_pred_gb)
```

```
plt.figure(figsize=(5,4))
sns.heatmap(cm, annot=True, fmt="d", cmap="RdPu", cbar=False,
            xticklabels=["Predicted 0", "Predicted 1"],
            yticklabels=["Actual 0", "Actual 1"])
plt.title("Confusion Matrix – Gradient Boosting")
plt.xlabel("Predicted Label")
plt.ylabel("True Label")
plt.show()
```




```
In [107... # 1) Get feature importances from trained Gradient Boosting model
importances = gboost.feature_importances_

# 2) Put into a DataFrame with feature names
feat_imp = pd.DataFrame({
    "Feature": X_train.columns,
    "Importance": importances
}).sort_values(by="Importance", ascending=False)

print(feat_imp)

# 3) Plot feature importance (bar chart)
plt.figure(figsize=(8,6))
sns.barplot(x="Importance", y="Feature", data=feat_imp, palette="Reds_r")
plt.title("Feature Importance – Gradient Boosting", fontsize=14)
plt.xlabel("Importance Score")
plt.ylabel("Feature")
plt.tight_layout()
plt.show()
```

	Feature	Importance
14	ST_Slope_Up	0.513421
3	Cholesterol	0.084141
7	Oldpeak	0.082765
6	ExerciseAngina	0.065509
5	MaxHR	0.056675
1	Sex	0.047185
8	ChestPainType_ATA	0.034654
9	ChestPainType_NAP	0.034181
2	RestingBP	0.025429
4	FastingBS	0.018967
10	ChestPainType_TA	0.013954
0	Age	0.012171
13	ST_Slope_Flat	0.007031
11	RestingECG_Normal	0.003777
12	RestingECG_ST	0.000138



NAIVE BAYES (GAUSSIAN)

```
In [110... # 1) Train
gnb = GaussianNB()
gnb.fit(X_train, y_train)

# 2) Predict
y_pred_nb = gnb.predict(X_test)
y_proba_nb = gnb.predict_proba(X_test)[:, 1]

# 3) Metrics (keep report & metrics, drop confusion matrix numbers)
```

```

nbAcc = accuracy_score(y_test, y_pred_nb)
nbAUC = roc_auc_score(y_test, y_proba_nb)

print("Naive Bayes (Gaussian) Accuracy:", nbAcc)
print("Naive Bayes (Gaussian) ROC AUC:", nbAUC)
print("\nClassification Report (GaussianNB):\n", classification_report(y_test, y_pred_nb))

# 4) ROC curve (brown tones)
fpr, tpr, _ = roc_curve(y_test, y_proba_nb)
plt.figure(figsize=(6,6))
plt.plot(fpr, tpr, lw=2, color="saddlebrown",
         label=f"GaussianNB (AUC = {nbAUC:.2f})")
plt.plot([0,1], [0,1], '--', color='lightgray')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("ROC Curve – Naive Bayes")
plt.legend(loc="lower right")
plt.show()

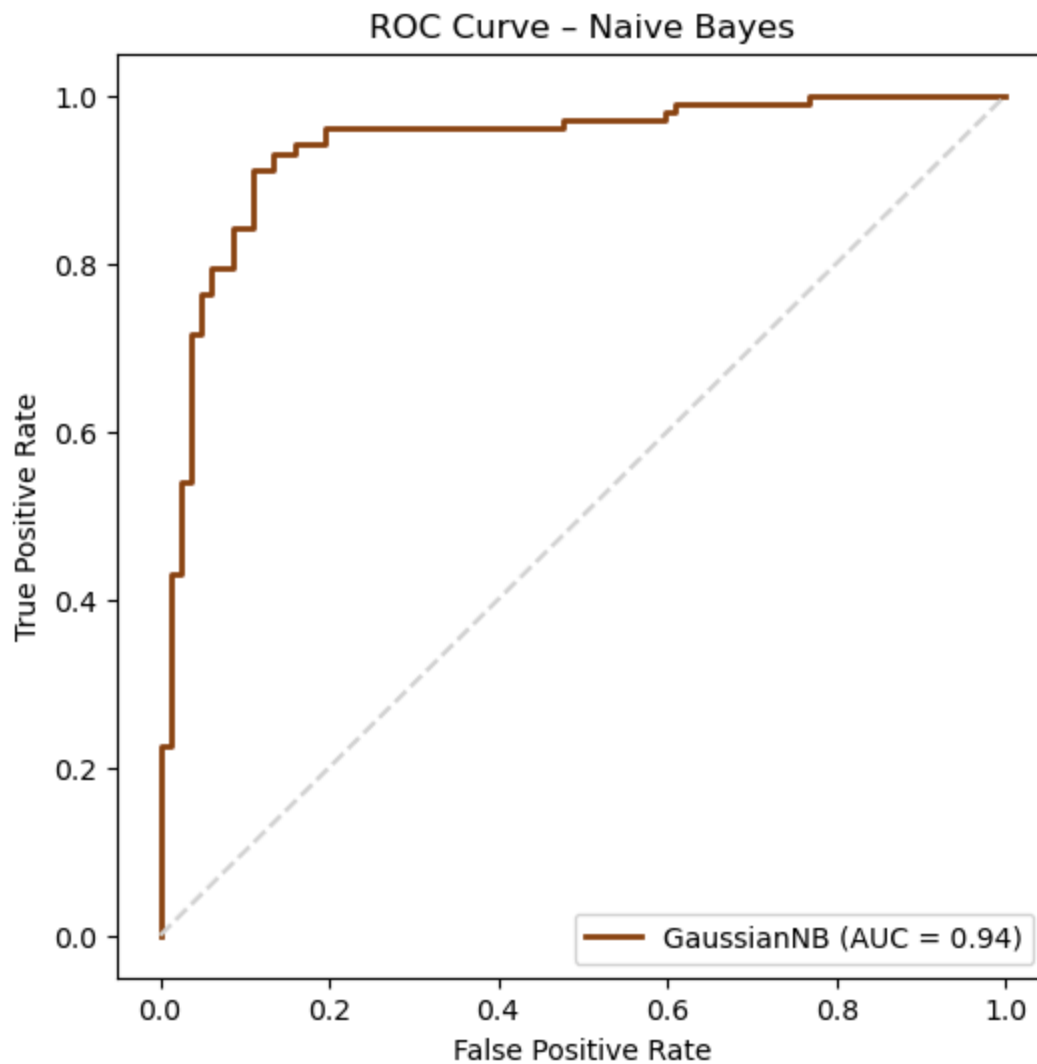
```

Naive Bayes (Gaussian) Accuracy: 0.8913043478260869

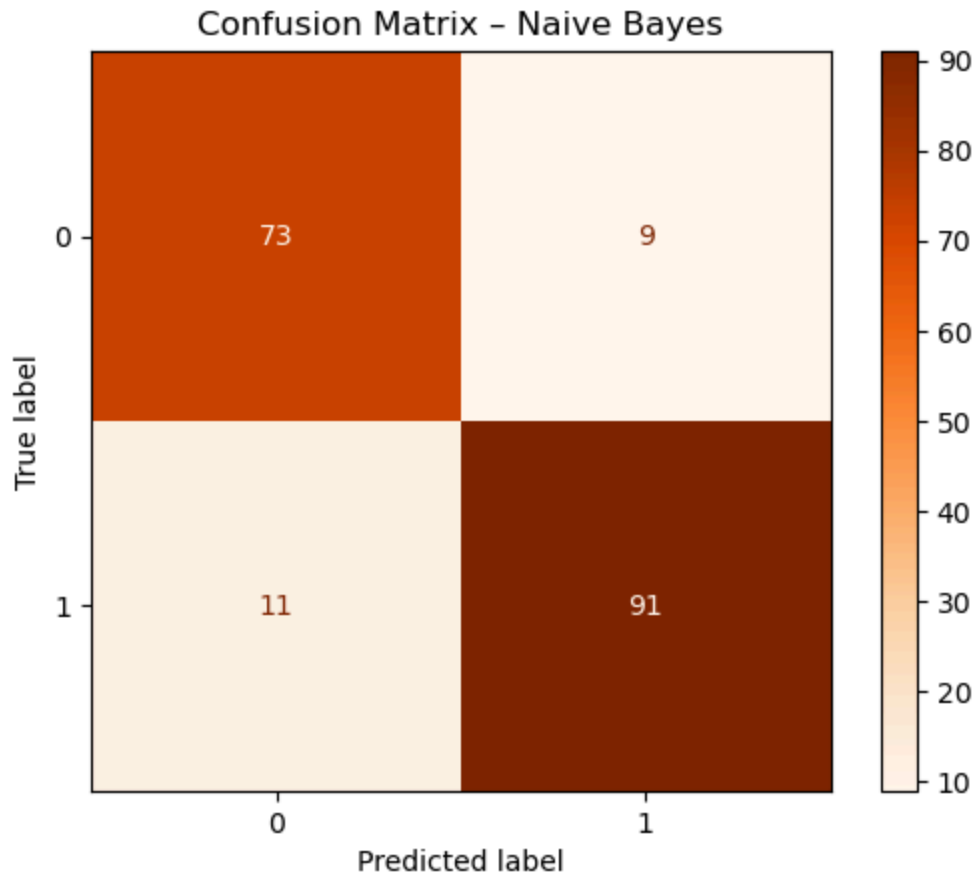
Naive Bayes (Gaussian) ROC AUC: 0.9404591104734578

Classification Report (GaussianNB):

	precision	recall	f1-score	support
0	0.87	0.89	0.88	82
1	0.91	0.89	0.90	102
accuracy			0.89	184
macro avg	0.89	0.89	0.89	184
weighted avg	0.89	0.89	0.89	184



```
In [112... # Confusion Matrix
cm = confusion_matrix(y_test, y_pred_nb)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap="Oranges")
plt.title("Confusion Matrix - Naive Bayes")
plt.show()
```

```
In [114... # Step 1: Create a dictionary with model performance results
data = {
    "Model": [
        "Logistic Regression",
        "Random Forest",
        "KNN (k=10)",
        "SVM (RBF)",
        "Decision Tree",
        "Gradient Boosting",
        "Naive Bayes"
    ],
    "Accuracy": [0.89, 0.88, 0.87, 0.89, 0.77, 0.91, 0.89],
    "Precision": [0.89, 0.88, 0.91, 0.89, 0.79, 0.93, 0.91],
    "Recall": [0.90, 0.90, 0.85, 0.90, 0.78, 0.90, 0.89],
    "F1-score": [0.90, 0.89, 0.88, 0.90, 0.78, 0.92, 0.90],
    "ROC AUC": [0.94, 0.94, 0.93, 0.94, 0.76, 0.95, 0.94]
}

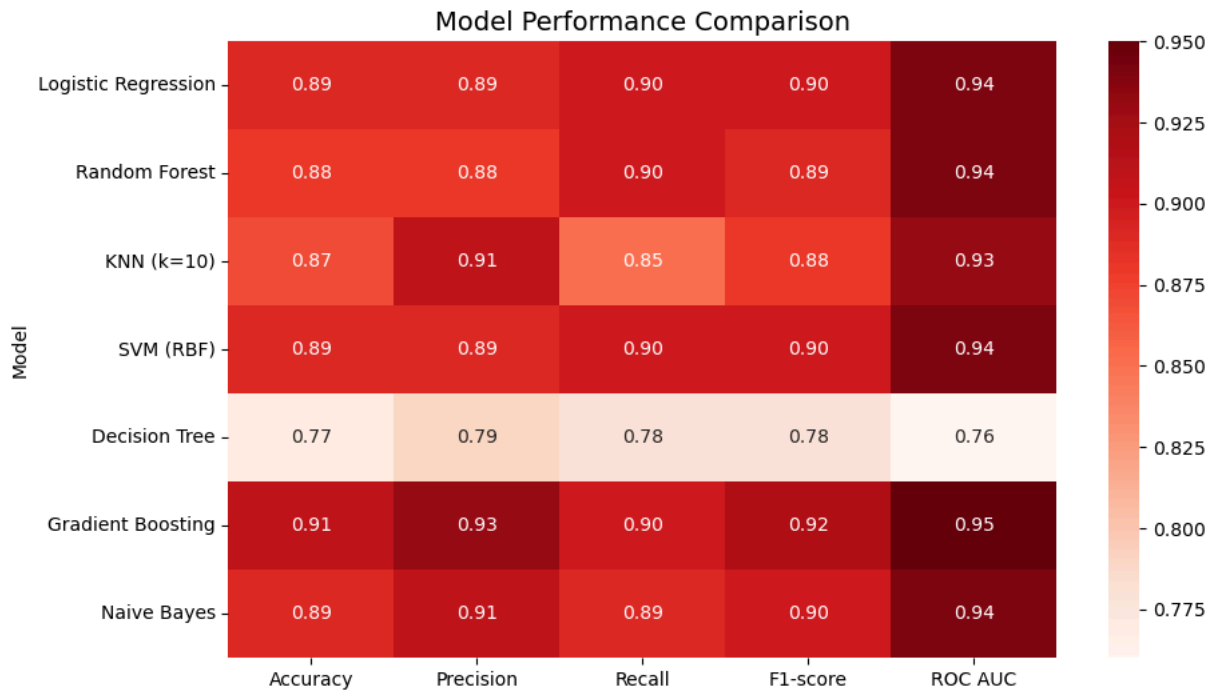
# Step 2: Convert to DataFrame
df_perf = pd.DataFrame(data)

# Step 3: Print table
print(df_perf)

# Step 4: Heatmap visualization
plt.figure(figsize=(10, 6))
sns.heatmap(df_perf.set_index("Model"), annot=True, cmap="Reds", fmt=".2f",
plt.title("Model Performance Comparison", fontsize=14)
```

```
plt.yticks(rotation=0)
plt.show()
```

	Model	Accuracy	Precision	Recall	F1-score	ROC AUC
0	Logistic Regression	0.89	0.89	0.90	0.90	0.94
1	Random Forest	0.88	0.88	0.90	0.89	0.94
2	KNN (k=10)	0.87	0.91	0.85	0.88	0.93
3	SVM (RBF)	0.89	0.89	0.90	0.90	0.94
4	Decision Tree	0.77	0.79	0.78	0.78	0.76
5	Gradient Boosting	0.91	0.93	0.90	0.92	0.95
6	Naive Bayes	0.89	0.91	0.89	0.90	0.94



Deep Learning

BASELINE MLP

```
In [118... # (only Dense layers, no regularization or callbacks)

# (Optional) set seeds for reproducibility
np.random.seed(42)
tf.random.set_seed(42)

# 1) Model architecture (very simple)
model_baseline = tf.keras.Sequential([
    tf.keras.layers.Dense(32, activation="relu", input_shape=(X_train.shape[1],)),
    tf.keras.layers.Dense(16, activation="relu"),
    tf.keras.layers.Dense(1, activation="sigmoid") # output layer for binary
])

# 2) Compile the model
# - Optimizer: Adam
# - Loss: Binary crossentropy (since target is 0/1)
# - Metrics: Accuracy and AUC
```

```

model_baseline.compile(
    optimizer="adam",
    loss="binary_crossentropy",
    metrics=["accuracy", tf.keras.metrics.AUC(name="auc")]
)

# 3) Training
# - Uses 20% of the training set for validation (validation_split=0.2)
# - Trains for 100 epochs with batch size = 32
# - No early stopping or learning rate scheduling
history_base = model_baseline.fit(
    X_train, y_train,
    epochs=100,
    batch_size=32,
    validation_split=0.2,
    verbose=1
)

# 4) Evaluation on the test set
# - Predictions are probabilities, converted to class labels at 0.5 threshold
proba_base = model_baseline.predict(X_test).ravel()
y_pred_base = (proba_base >= 0.5).astype(int)

print("Test Accuracy (Baseline MLP):", accuracy_score(y_test, y_pred_base))
print("Test AUC (Baseline MLP):", roc_auc_score(y_test, proba_base))
print("\nClassification Report (Baseline MLP):\n", classification_report(y_test, y_pred_base))
print("Confusion Matrix (Baseline MLP):\n", confusion_matrix(y_test, y_pred_base))

```




















Epoch 1/100

```


2025-08-29 17:36:05.921806: I metal_plugin/src/device/metal_device.cc:1154]
Metal device set to: Apple M3 Pro
2025-08-29 17:36:05.921903: I metal_plugin/src/device/metal_device.cc:296] s
ystemMemory: 18.00 GB
2025-08-29 17:36:05.921924: I metal_plugin/src/device/metal_device.cc:313] m
axCacheSize: 6.00 GB
2025-08-29 17:36:05.921975: I tensorflow/core/common_runtime/pluggable_devic
e/pluggable_device_factory.cc:305] Could not identify NUMA node of platform
GPU ID 0, defaulting to 0. Your kernel may not have been built with NUMA sup
port.
2025-08-29 17:36:05.922011: I tensorflow/core/common_runtime/pluggable_devic
e/pluggable_device_factory.cc:271] Created TensorFlow device (/job:localhos
t/replica:0/task:0/device:GPU:0 with 0 MB memory) -> physical PluggableDevic
e (device: 0, name: METAL, pci bus id: <undefined>)
2025-08-29 17:36:06.229703: I tensorflow/core/grappler/optimizers/custom_gra
ph_optimizer_registry.cc:117] Plugin optimizer for device_type GPU is enable
d.
2025-08-29 17:36:06.230842: E tensorflow/core/grappler/optimizers/meta_optim
izer.cc:961] PluggableGraphOptimizer failed: INVALID_ARGUMENT: Failed to des
erialize the `graph_buf`.


```


```


19/19  2s 41ms/step - accuracy: 0.5853 - auc: 0.7332 - loss: 0.6408 - val_accuracy: 0.6871 - val_auc: 0.8036 - val_loss: 0.6008
Epoch 2/100
19/19  0s 10ms/step - accuracy: 0.7474 - auc: 0.8732 - loss: 0.5550 - val_accuracy: 0.7551 - val_auc: 0.8416 - val_loss: 0.5469
Epoch 3/100
19/19  0s 10ms/step - accuracy: 0.8242 - auc: 0.8971 - loss: 0.4898 - val_accuracy: 0.7755 - val_auc: 0.8423 - val_loss: 0.5074
Epoch 4/100
19/19  0s 10ms/step - accuracy: 0.8464 - auc: 0.9060 - loss: 0.4382 - val_accuracy: 0.8027 - val_auc: 0.8471 - val_loss: 0.4823
Epoch 5/100
19/19  0s 10ms/step - accuracy: 0.8567 - auc: 0.9121 - loss: 0.4011 - val_accuracy: 0.8095 - val_auc: 0.8495 - val_loss: 0.4678
Epoch 6/100
19/19  0s 10ms/step - accuracy: 0.8618 - auc: 0.9171 - loss: 0.3761 - val_accuracy: 0.8095 - val_auc: 0.8537 - val_loss: 0.4601
Epoch 7/100
19/19  0s 10ms/step - accuracy: 0.8635 - auc: 0.9216 - loss: 0.3592 - val_accuracy: 0.8027 - val_auc: 0.8557 - val_loss: 0.4561
Epoch 8/100
19/19  0s 10ms/step - accuracy: 0.8635 - auc: 0.9264 - loss: 0.3465 - val_accuracy: 0.8027 - val_auc: 0.8575 - val_loss: 0.4541
Epoch 9/100
19/19  0s 10ms/step - accuracy: 0.8720 - auc: 0.9300 - loss: 0.3362 - val_accuracy: 0.8095 - val_auc: 0.8603 - val_loss: 0.4528
Epoch 10/100
19/19  0s 10ms/step - accuracy: 0.8720 - auc: 0.9334 - loss: 0.3278 - val_accuracy: 0.8095 - val_auc: 0.8611 - val_loss: 0.4520
Epoch 11/100
19/19  0s 10ms/step - accuracy: 0.8720 - auc: 0.9365 - loss: 0.3204 - val_accuracy: 0.8095 - val_auc: 0.8605 - val_loss: 0.4515
Epoch 12/100
19/19  0s 10ms/step - accuracy: 0.8754 - auc: 0.9392 - loss: 0.3136 - val_accuracy: 0.8027 - val_auc: 0.8620 - val_loss: 0.4518
Epoch 13/100
19/19  0s 10ms/step - accuracy: 0.8754 - auc: 0.9417 - loss: 0.3075 - val_accuracy: 0.8027 - val_auc: 0.8616 - val_loss: 0.4527
Epoch 14/100
19/19  0s 10ms/step - accuracy: 0.8771 - auc: 0.9436 - loss: 0.3019 - val_accuracy: 0.8095 - val_auc: 0.8620 - val_loss: 0.4538
Epoch 15/100
19/19  0s 10ms/step - accuracy: 0.8805 - auc: 0.9457 - loss: 0.2968 - val_accuracy: 0.8163 - val_auc: 0.8624 - val_loss: 0.4550
Epoch 16/100
19/19  0s 10ms/step - accuracy: 0.8823 - auc: 0.9474 - loss: 0.2922 - val_accuracy: 0.8231 - val_auc: 0.8618 - val_loss: 0.4560
Epoch 17/100
19/19  0s 10ms/step - accuracy: 0.8840 - auc: 0.9491 - loss: 0.2876 - val_accuracy: 0.8299 - val_auc: 0.8624 - val_loss: 0.4570
Epoch 18/100
19/19  0s 10ms/step - accuracy: 0.8857 - auc: 0.9509 - loss: 0.2834 - val_accuracy: 0.8299 - val_auc: 0.8630 - val_loss: 0.4580
Epoch 19/100
19/19  0s 10ms/step - accuracy: 0.8891 - auc: 0.9521 - loss: 0.2795 - val_accuracy: 0.8299 - val_auc: 0.8626 - val_loss: 0.4593


```


Epoch 20/100
19/19  0s 10ms/step - accuracy: 0.8874 - auc: 0.9532 - loss: 0.2760 - val_accuracy: 0.8299 - val_auc: 0.8635 - val_loss: 0.4602


Epoch 21/100
19/19  0s 10ms/step - accuracy: 0.8908 - auc: 0.9544 - loss: 0.2727 - val_accuracy: 0.8299 - val_auc: 0.8642 - val_loss: 0.4614


Epoch 22/100
19/19  0s 10ms/step - accuracy: 0.8908 - auc: 0.9557 - loss: 0.2696 - val_accuracy: 0.8231 - val_auc: 0.8655 - val_loss: 0.4628


Epoch 23/100
19/19  0s 10ms/step - accuracy: 0.8959 - auc: 0.9568 - loss: 0.2667 - val_accuracy: 0.8231 - val_auc: 0.8657 - val_loss: 0.4639

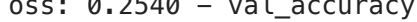
Epoch 24/100
19/19  0s 10ms/step - accuracy: 0.8959 - auc: 0.9577 - loss: 0.2639 - val_accuracy: 0.8231 - val_auc: 0.8654 - val_loss: 0.4654

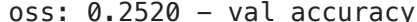
Epoch 25/100
19/19  0s 10ms/step - accuracy: 0.8976 - auc: 0.9584 - loss: 0.2614 - val_accuracy: 0.8299 - val_auc: 0.8662 - val_loss: 0.4672


Epoch 26/100
19/19  0s 10ms/step - accuracy: 0.8976 - auc: 0.9592 - loss: 0.2587 - val_accuracy: 0.8299 - val_auc: 0.8670 - val_loss: 0.4689


Epoch 27/100
19/19  0s 10ms/step - accuracy: 0.8959 - auc: 0.9601 - loss: 0.2564 - val_accuracy: 0.8299 - val_auc: 0.8675 - val_loss: 0.4703


Epoch 28/100
19/19  0s 10ms/step - accuracy: 0.8942 - auc: 0.9607 - loss: 0.2540 - val_accuracy: 0.8299 - val_auc: 0.8661 - val_loss: 0.4725


Epoch 29/100
19/19  0s 10ms/step - accuracy: 0.8942 - auc: 0.9612 - loss: 0.2520 - val_accuracy: 0.8163 - val_auc: 0.8655 - val_loss: 0.4740


Epoch 30/100
19/19  0s 10ms/step - accuracy: 0.8959 - auc: 0.9620 - loss: 0.2497 - val_accuracy: 0.8163 - val_auc: 0.8656 - val_loss: 0.4753


Epoch 31/100
19/19  0s 10ms/step - accuracy: 0.8959 - auc: 0.9626 - loss: 0.2477 - val_accuracy: 0.8095 - val_auc: 0.8663 - val_loss: 0.4775


Epoch 32/100
19/19  0s 10ms/step - accuracy: 0.8976 - auc: 0.9630 - loss: 0.2459 - val_accuracy: 0.8095 - val_auc: 0.8664 - val_loss: 0.4785


Epoch 33/100
19/19  0s 11ms/step - accuracy: 0.8976 - auc: 0.9636 - loss: 0.2438 - val_accuracy: 0.8095 - val_auc: 0.8667 - val_loss: 0.4801



















Epoch 34/100
19/19  0s 13ms/step - accuracy: 0.8976 - auc: 0.9642 - loss: 0.2420 - val_accuracy: 0.8095 - val_auc: 0.8655 - val_loss: 0.4811


Epoch 35/100
19/19  0s 11ms/step - accuracy: 0.8993 - auc: 0.9649 - loss: 0.2400 - val_accuracy: 0.8095 - val_auc: 0.8661 - val_loss: 0.4827


Epoch 36/100
19/19  0s 11ms/step - accuracy: 0.9010 - auc: 0.9655 - loss: 0.2383 - val_accuracy: 0.8163 - val_auc: 0.8670 - val_loss: 0.4836


Epoch 37/100
19/19  0s 11ms/step - accuracy: 0.8993 - auc: 0.9661 - loss: 0.2364 - val_accuracy: 0.8163 - val_auc: 0.8663 - val_loss: 0.4847


Epoch 38/100
19/19  0s 12ms/step - accuracy: 0.8993 - auc: 0.9667 - loss: 0.2345 - val_accuracy: 0.8163 - val_auc: 0.8663 - val_loss: 0.4847


oss: 0.2345 - val_accuracy: 0.8163 - val_auc: 0.8666 - val_loss: 0.4861
Epoch 39/100
19/19  0s 10ms/step - accuracy: 0.9027 - auc: 0.9671 - l
oss: 0.2329 - val_accuracy: 0.8027 - val_auc: 0.8660 - val_loss: 0.4865
Epoch 40/100
19/19  0s 10ms/step - accuracy: 0.9044 - auc: 0.9678 - l
oss: 0.2308 - val_accuracy: 0.8027 - val_auc: 0.8662 - val_loss: 0.4880
Epoch 41/100
19/19  0s 12ms/step - accuracy: 0.9044 - auc: 0.9682 - l
oss: 0.2294 - val_accuracy: 0.8027 - val_auc: 0.8670 - val_loss: 0.4890
Epoch 42/100
19/19  0s 10ms/step - accuracy: 0.9027 - auc: 0.9686 - l
oss: 0.2276 - val_accuracy: 0.8027 - val_auc: 0.8666 - val_loss: 0.4905
Epoch 43/100
19/19  0s 10ms/step - accuracy: 0.9044 - auc: 0.9690 - l
oss: 0.2263 - val_accuracy: 0.8027 - val_auc: 0.8670 - val_loss: 0.4916
Epoch 44/100
19/19  0s 10ms/step - accuracy: 0.9061 - auc: 0.9695 - l
oss: 0.2246 - val_accuracy: 0.8027 - val_auc: 0.8649 - val_loss: 0.4926
Epoch 45/100
19/19  0s 12ms/step - accuracy: 0.9078 - auc: 0.9699 - l
oss: 0.2232 - val_accuracy: 0.8027 - val_auc: 0.8658 - val_loss: 0.4939
Epoch 46/100
19/19  0s 11ms/step - accuracy: 0.9078 - auc: 0.9704 - l
oss: 0.2216 - val_accuracy: 0.8027 - val_auc: 0.8661 - val_loss: 0.4956
Epoch 47/100
19/19  0s 11ms/step - accuracy: 0.9078 - auc: 0.9705 - l
oss: 0.2204 - val_accuracy: 0.8027 - val_auc: 0.8665 - val_loss: 0.4959
Epoch 48/100
19/19  0s 10ms/step - accuracy: 0.9078 - auc: 0.9710 - l
oss: 0.2188 - val_accuracy: 0.8027 - val_auc: 0.8672 - val_loss: 0.4974
Epoch 49/100
19/19  0s 11ms/step - accuracy: 0.9130 - auc: 0.9713 - l
oss: 0.2175 - val_accuracy: 0.7891 - val_auc: 0.8667 - val_loss: 0.4983
Epoch 50/100
19/19  0s 11ms/step - accuracy: 0.9113 - auc: 0.9717 - l
oss: 0.2161 - val_accuracy: 0.7891 - val_auc: 0.8674 - val_loss: 0.4996
Epoch 51/100
19/19  0s 11ms/step - accuracy: 0.9147 - auc: 0.9722 - l
oss: 0.2147 - val_accuracy: 0.7891 - val_auc: 0.8674 - val_loss: 0.5011
Epoch 52/100
19/19  0s 11ms/step - accuracy: 0.9147 - auc: 0.9725 - l
oss: 0.2136 - val_accuracy: 0.7891 - val_auc: 0.8680 - val_loss: 0.5021
Epoch 53/100
19/19  0s 10ms/step - accuracy: 0.9164 - auc: 0.9729 - l
oss: 0.2122 - val_accuracy: 0.7891 - val_auc: 0.8689 - val_loss: 0.5030
Epoch 54/100
19/19  0s 12ms/step - accuracy: 0.9164 - auc: 0.9731 - l
oss: 0.2110 - val_accuracy: 0.7891 - val_auc: 0.8688 - val_loss: 0.5042
Epoch 55/100
19/19  0s 11ms/step - accuracy: 0.9181 - auc: 0.9733 - l
oss: 0.2097 - val_accuracy: 0.7891 - val_auc: 0.8694 - val_loss: 0.5048
Epoch 56/100
19/19  0s 11ms/step - accuracy: 0.9198 - auc: 0.9736 - l
oss: 0.2085 - val_accuracy: 0.7891 - val_auc: 0.8697 - val_loss: 0.5054
Epoch 57/100


19/19  0s 11ms/step - accuracy: 0.9198 - auc: 0.9738 - loss: 0.2074 - val_accuracy: 0.7891 - val_auc: 0.8695 - val_loss: 0.5065
Epoch 58/100


19/19  0s 12ms/step - accuracy: 0.9198 - auc: 0.9742 - loss: 0.2059 - val_accuracy: 0.7891 - val_auc: 0.8684 - val_loss: 0.5068
Epoch 59/100


19/19  0s 10ms/step - accuracy: 0.9215 - auc: 0.9745 - loss: 0.2050 - val_accuracy: 0.7891 - val_auc: 0.8694 - val_loss: 0.5076
Epoch 60/100


19/19  0s 11ms/step - accuracy: 0.9198 - auc: 0.9748 - loss: 0.2036 - val_accuracy: 0.7891 - val_auc: 0.8695 - val_loss: 0.5089
Epoch 61/100


19/19  0s 10ms/step - accuracy: 0.9215 - auc: 0.9751 - loss: 0.2025 - val_accuracy: 0.7891 - val_auc: 0.8694 - val_loss: 0.5086
Epoch 62/100


19/19  0s 12ms/step - accuracy: 0.9232 - auc: 0.9753 - loss: 0.2013 - val_accuracy: 0.7891 - val_auc: 0.8677 - val_loss: 0.5104
Epoch 63/100

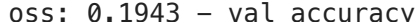
19/19  0s 11ms/step - accuracy: 0.9249 - auc: 0.9755 - loss: 0.2001 - val_accuracy: 0.7891 - val_auc: 0.8674 - val_loss: 0.5100
Epoch 64/100


19/19  0s 10ms/step - accuracy: 0.9266 - auc: 0.9757 - loss: 0.1990 - val_accuracy: 0.7891 - val_auc: 0.8675 - val_loss: 0.5118
Epoch 65/100


19/19  0s 10ms/step - accuracy: 0.9283 - auc: 0.9760 - loss: 0.1979 - val_accuracy: 0.7891 - val_auc: 0.8675 - val_loss: 0.5110
Epoch 66/100


19/19  0s 10ms/step - accuracy: 0.9266 - auc: 0.9767 - loss: 0.1966 - val_accuracy: 0.7959 - val_auc: 0.8676 - val_loss: 0.5125
Epoch 67/100


19/19  0s 10ms/step - accuracy: 0.9283 - auc: 0.9769 - loss: 0.1955 - val_accuracy: 0.7959 - val_auc: 0.8678 - val_loss: 0.5126
Epoch 68/100


19/19  0s 10ms/step - accuracy: 0.9266 - auc: 0.9772 - loss: 0.1943 - val_accuracy: 0.7959 - val_auc: 0.8680 - val_loss: 0.5132
Epoch 69/100


19/19  0s 10ms/step - accuracy: 0.9283 - auc: 0.9775 - loss: 0.1932 - val_accuracy: 0.7959 - val_auc: 0.8684 - val_loss: 0.5140
Epoch 70/100


19/19  0s 10ms/step - accuracy: 0.9283 - auc: 0.9778 - loss: 0.1922 - val_accuracy: 0.7959 - val_auc: 0.8683 - val_loss: 0.5141
Epoch 71/100


19/19  0s 12ms/step - accuracy: 0.9283 - auc: 0.9780 - loss: 0.1911 - val_accuracy: 0.7959 - val_auc: 0.8684 - val_loss: 0.5160
Epoch 72/100


19/19  0s 12ms/step - accuracy: 0.9300 - auc: 0.9783 - loss: 0.1899 - val_accuracy: 0.7959 - val_auc: 0.8689 - val_loss: 0.5162
Epoch 73/100


19/19  0s 11ms/step - accuracy: 0.9300 - auc: 0.9786 - loss: 0.1890 - val_accuracy: 0.7891 - val_auc: 0.8673 - val_loss: 0.5176
Epoch 74/100


19/19  0s 12ms/step - accuracy: 0.9300 - auc: 0.9788 - loss: 0.1878 - val_accuracy: 0.7891 - val_auc: 0.8679 - val_loss: 0.5181
Epoch 75/100


19/19  0s 11ms/step - accuracy: 0.9300 - auc: 0.9791 - loss: 0.1869 - val_accuracy: 0.7891 - val_auc: 0.8675 - val_loss: 0.5187


Epoch 76/100
19/19  0s 11ms/step - accuracy: 0.9300 - auc: 0.9793 - loss: 0.1857 - val_accuracy: 0.7891 - val_auc: 0.8676 - val_loss: 0.5199


Epoch 77/100
19/19  0s 10ms/step - accuracy: 0.9300 - auc: 0.9795 - loss: 0.1847 - val_accuracy: 0.7891 - val_auc: 0.8687 - val_loss: 0.5211


Epoch 78/100
19/19  0s 10ms/step - accuracy: 0.9300 - auc: 0.9797 - loss: 0.1836 - val_accuracy: 0.7891 - val_auc: 0.8690 - val_loss: 0.5219


Epoch 79/100
19/19  0s 10ms/step - accuracy: 0.9300 - auc: 0.9799 - loss: 0.1826 - val_accuracy: 0.7891 - val_auc: 0.8687 - val_loss: 0.5223

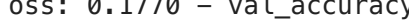
Epoch 80/100
19/19  0s 10ms/step - accuracy: 0.9300 - auc: 0.9803 - loss: 0.1813 - val_accuracy: 0.7891 - val_auc: 0.8680 - val_loss: 0.5244

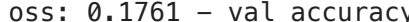
Epoch 81/100
19/19  0s 9ms/step - accuracy: 0.9283 - auc: 0.9805 - loss: 0.1806 - val_accuracy: 0.7891 - val_auc: 0.8680 - val_loss: 0.5245


Epoch 82/100
19/19  0s 10ms/step - accuracy: 0.9300 - auc: 0.9810 - loss: 0.1792 - val_accuracy: 0.7891 - val_auc: 0.8685 - val_loss: 0.5261


Epoch 83/100
19/19  0s 10ms/step - accuracy: 0.9283 - auc: 0.9811 - loss: 0.1782 - val_accuracy: 0.7891 - val_auc: 0.8673 - val_loss: 0.5270


Epoch 84/100
19/19  0s 10ms/step - accuracy: 0.9317 - auc: 0.9813 - loss: 0.1770 - val_accuracy: 0.7891 - val_auc: 0.8671 - val_loss: 0.5287


Epoch 85/100
19/19  0s 10ms/step - accuracy: 0.9317 - auc: 0.9815 - loss: 0.1761 - val_accuracy: 0.7891 - val_auc: 0.8674 - val_loss: 0.5293


Epoch 86/100
19/19  0s 9ms/step - accuracy: 0.9317 - auc: 0.9818 - loss: 0.1747 - val_accuracy: 0.7891 - val_auc: 0.8673 - val_loss: 0.5311


Epoch 87/100
19/19  0s 11ms/step - accuracy: 0.9300 - auc: 0.9821 - loss: 0.1738 - val_accuracy: 0.7891 - val_auc: 0.8673 - val_loss: 0.5314


Epoch 88/100
19/19  0s 10ms/step - accuracy: 0.9334 - auc: 0.9824 - loss: 0.1724 - val_accuracy: 0.7823 - val_auc: 0.8664 - val_loss: 0.5338


Epoch 89/100
19/19  0s 10ms/step - accuracy: 0.9352 - auc: 0.9827 - loss: 0.1713 - val_accuracy: 0.7891 - val_auc: 0.8648 - val_loss: 0.5348

Epoch 90/100
19/19  0s 10ms/step - accuracy: 0.9352 - auc: 0.9828 - loss: 0.1701 - val_accuracy: 0.7823 - val_auc: 0.8656 - val_loss: 0.5362

Epoch 91/100
19/19  0s 10ms/step - accuracy: 0.9352 - auc: 0.9831 - loss: 0.1691 - val_accuracy: 0.7823 - val_auc: 0.8652 - val_loss: 0.5370

Epoch 92/100
19/19  0s 10ms/step - accuracy: 0.9369 - auc: 0.9834 - loss: 0.1676 - val_accuracy: 0.7823 - val_auc: 0.8648 - val_loss: 0.5388

Epoch 93/100
19/19  0s 10ms/step - accuracy: 0.9369 - auc: 0.9837 - loss: 0.1667 - val_accuracy: 0.7823 - val_auc: 0.8644 - val_loss: 0.5403

Epoch 94/100
19/19  0s 12ms/step - accuracy: 0.9403 - auc: 0.9839 - loss: 0.1657 - val_accuracy: 0.7823 - val_auc: 0.8644 - val_loss: 0.5403


```

loss: 0.1655 - val_accuracy: 0.7823 - val_auc: 0.8628 - val_loss: 0.5414
Epoch 95/100
19/19 ━━━━━━━━━━━━━━━━━━━ 0s 10ms/step - accuracy: 0.9403 - auc: 0.9842 - l
loss: 0.1642 - val_accuracy: 0.7823 - val_auc: 0.8633 - val_loss: 0.5424
Epoch 96/100
19/19 ━━━━━━━━━━━━━━━━━━━ 0s 10ms/step - accuracy: 0.9386 - auc: 0.9844 - l
loss: 0.1633 - val_accuracy: 0.7823 - val_auc: 0.8630 - val_loss: 0.5428
Epoch 97/100
19/19 ━━━━━━━━━━━━━━━━━━━ 0s 10ms/step - accuracy: 0.9403 - auc: 0.9848 - l
loss: 0.1619 - val_accuracy: 0.7823 - val_auc: 0.8630 - val_loss: 0.5452
Epoch 98/100
19/19 ━━━━━━━━━━━━━━━━━━━ 0s 10ms/step - accuracy: 0.9403 - auc: 0.9850 - l
loss: 0.1608 - val_accuracy: 0.7823 - val_auc: 0.8628 - val_loss: 0.5460
Epoch 99/100
19/19 ━━━━━━━━━━━━━━━━━━━ 0s 10ms/step - accuracy: 0.9403 - auc: 0.9853 - l
loss: 0.1597 - val_accuracy: 0.7823 - val_auc: 0.8629 - val_loss: 0.5470
Epoch 100/100
19/19 ━━━━━━━━━━━━━━━━━━━ 0s 9ms/step - accuracy: 0.9403 - auc: 0.9854 - lo
ss: 0.1584 - val_accuracy: 0.7823 - val_auc: 0.8632 - val_loss: 0.5495
6/6 ━━━━━━━━━━━━━━━━━━━ 0s 6ms/step
Test Accuracy (Baseline MLP): 0.875
Test AUC (Baseline MLP): 0.9230033476805356

```

Classification Report (Baseline MLP):

	precision	recall	f1-score	support
0	0.85	0.88	0.86	82
1	0.90	0.87	0.89	102
accuracy			0.88	184
macro avg	0.87	0.88	0.87	184
weighted avg	0.88	0.88	0.88	184

Confusion Matrix (Baseline MLP):

```
[[72 10]
 [13 89]]
```

Enhanced MLP

```

In [120... # (Optional) ensure reproducibility
np.random.seed(42)
tf.random.set_seed(42)

# 1) Define a simple MLP architecture for binary classification
# - Hidden layers use ReLU to capture non-linearities
# - BatchNorm + Dropout help stabilize training and reduce overfitting
# - Output layer is a single neuron with Sigmoid to output P(y=1)
model = Sequential([
    Dense(64, activation="relu", input_shape=(X_train.shape[1],)),
    BatchNormalization(),
    Dropout(0.2),

    Dense(32, activation="relu"),
    BatchNormalization(),
    Dropout(0.2),

```

```

        Dense(16, activation="relu"),
        Dense(1, activation="sigmoid") # binary classification output (probability)
    ])

# 2) Compile the model
# - Binary cross-entropy is the standard loss for 0/1 targets
# - Track both Accuracy and AUC (discrimination)
model.compile(
    optimizer="adam",
    loss="binary_crossentropy",
    metrics=["accuracy", tf.keras.metrics.AUC(name="auc")]
)

# 3) Callbacks to improve training stability
# - EarlyStopping: stop when val_loss stops improving; restore best weights
# - ReduceLROnPlateau: reduce learning rate when val_loss plateaus
early = EarlyStopping(monitor="val_loss", patience=10, restore_best_weights=True)
plateau = ReduceLROnPlateau(monitor="val_loss", factor=0.5, patience=5, min_lr=1e-6)


# 4) Train the model
# - validation_split=0.2: uses 20% of the TRAINING set as validation
#   (test set remains untouched for final evaluation)
history = model.fit(
    X_train, y_train,
    epochs=100,
    batch_size=32,
    validation_split=0.2,
    callbacks=[early, plateau],
    verbose=1
)


# 5) Final evaluation on the held-out test set
# - proba: predicted probability of the positive class (y=1)
# - y_pred: hard labels using a default 0.50 threshold
proba = model.predict(X_test).ravel()
y_pred = (proba >= 0.5).astype(int)


print("Test Accuracy:", accuracy_score(y_test, y_pred))
print("Test AUC:", roc_auc_score(y_test, proba))
print("\nClassification Report:\n", classification_report(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))


# (Optional) If you want to optimize the decision threshold instead of using 0.5
# you can compute the ROC curve and pick the Youden J point (argmax of TPR - FPR)
# Then recompute y_pred and the confusion matrix at that threshold.


```


Epoch 1/100
19/19  2s 46ms/step - accuracy: 0.5512 - auc: 0.6266 - loss: 0.7772 - val_accuracy: 0.6599 - val_auc: 0.7215 - val_loss: 0.6653 - learning_rate: 0.0010


Epoch 2/100
19/19  0s 18ms/step - accuracy: 0.7014 - auc: 0.8002 - loss: 0.5732 - val_accuracy: 0.7211 - val_auc: 0.8085 - val_loss: 0.6049 - learning_rate: 0.0010


Epoch 3/100
19/19  0s 17ms/step - accuracy: 0.8020 - auc: 0.8790 - loss: 0.4532 - val_accuracy: 0.7755 - val_auc: 0.8312 - val_loss: 0.5646 - learning_rate: 0.0010


Epoch 4/100
19/19  0s 16ms/step - accuracy: 0.8225 - auc: 0.8936 - loss: 0.4173 - val_accuracy: 0.7959 - val_auc: 0.8528 - val_loss: 0.5278 - learning_rate: 0.0010


Epoch 5/100
19/19  0s 16ms/step - accuracy: 0.8311 - auc: 0.8974 - loss: 0.4064 - val_accuracy: 0.8095 - val_auc: 0.8538 - val_loss: 0.5042 - learning_rate: 0.0010


Epoch 6/100
19/19  0s 16ms/step - accuracy: 0.8345 - auc: 0.9137 - loss: 0.3667 - val_accuracy: 0.8095 - val_auc: 0.8570 - val_loss: 0.4880 - learning_rate: 0.0010


Epoch 7/100
19/19  0s 16ms/step - accuracy: 0.8447 - auc: 0.9195 - loss: 0.3561 - val_accuracy: 0.8231 - val_auc: 0.8553 - val_loss: 0.4746 - learning_rate: 0.0010


Epoch 8/100
19/19  0s 16ms/step - accuracy: 0.8737 - auc: 0.9324 - loss: 0.3298 - val_accuracy: 0.8299 - val_auc: 0.8565 - val_loss: 0.4642 - learning_rate: 0.0010


Epoch 9/100
19/19  0s 16ms/step - accuracy: 0.8532 - auc: 0.9257 - loss: 0.3442 - val_accuracy: 0.8299 - val_auc: 0.8570 - val_loss: 0.4580 - learning_rate: 0.0010

Epoch 10/100
19/19  0s 16ms/step - accuracy: 0.8601 - auc: 0.9412 - loss: 0.3087 - val_accuracy: 0.8367 - val_auc: 0.8604 - val_loss: 0.4537 - learning_rate: 0.0010


Epoch 11/100
19/19  0s 16ms/step - accuracy: 0.8652 - auc: 0.9350 - loss: 0.3181 - val_accuracy: 0.8367 - val_auc: 0.8615 - val_loss: 0.4498 - learning_rate: 0.0010

Epoch 12/100
19/19  0s 16ms/step - accuracy: 0.8652 - auc: 0.9486 - loss: 0.2906 - val_accuracy: 0.8435 - val_auc: 0.8615 - val_loss: 0.4499 - learning_rate: 0.0010


Epoch 13/100
19/19  0s 17ms/step - accuracy: 0.8669 - auc: 0.9435 - loss: 0.2999 - val_accuracy: 0.8367 - val_auc: 0.8581 - val_loss: 0.4521 - learning_rate: 0.0010

Epoch 14/100
19/19  0s 18ms/step - accuracy: 0.8908 - auc: 0.9509 - loss: 0.2788 - val_accuracy: 0.8299 - val_auc: 0.8582 - val_loss: 0.4520 - learning_rate: 0.0010


Epoch 15/100

19/19  0s 17ms/step - accuracy: 0.8840 - auc: 0.9496 - loss: 0.2866 - val_accuracy: 0.8231 - val_auc: 0.8585 - val_loss: 0.4519 - learning_rate: 0.0010


Epoch 16/100

19/19  0s 17ms/step - accuracy: 0.8874 - auc: 0.9570 - loss: 0.2683 - val_accuracy: 0.8163 - val_auc: 0.8573 - val_loss: 0.4541 - learning_rate: 0.0010


Epoch 17/100

19/19  0s 16ms/step - accuracy: 0.8908 - auc: 0.9560 - loss: 0.2662 - val_accuracy: 0.8163 - val_auc: 0.8563 - val_loss: 0.4553 - learning_rate: 5.0000e-04


Epoch 18/100

19/19  0s 16ms/step - accuracy: 0.8891 - auc: 0.9572 - loss: 0.2628 - val_accuracy: 0.8163 - val_auc: 0.8554 - val_loss: 0.4575 - learning_rate: 5.0000e-04


Epoch 19/100

19/19  0s 16ms/step - accuracy: 0.8976 - auc: 0.9473 - loss: 0.2857 - val_accuracy: 0.8027 - val_auc: 0.8583 - val_loss: 0.4600 - learning_rate: 5.0000e-04

Epoch 20/100

19/19  0s 17ms/step - accuracy: 0.8788 - auc: 0.9567 - loss: 0.2637 - val_accuracy: 0.8095 - val_auc: 0.8556 - val_loss: 0.4633 - learning_rate: 5.0000e-04

Epoch 21/100

19/19  0s 16ms/step - accuracy: 0.8976 - auc: 0.9623 - loss: 0.2463 - val_accuracy: 0.7959 - val_auc: 0.8578 - val_loss: 0.4649 - learning_rate: 5.0000e-04

6/6  0s 12ms/step

Test Accuracy: 0.8586956521739131

Test AUC: 0.9330463892874223

Classification Report:

	precision	recall	f1-score	support
0	0.92	0.74	0.82	82
1	0.82	0.95	0.88	102
accuracy			0.86	184
macro avg	0.87	0.85	0.85	184
weighted avg	0.87	0.86	0.86	184

Confusion Matrix:

```
[[61 21]
 [ 5 97]]
```

LSTM

```
In [142... # 1) Convert all features to numeric.
# - If a value cannot be converted, it will be set to NaN.
X_num = X.apply(pd.to_numeric, errors='coerce')

# 2) Handle missing values (NaN).
# - Here we replace them with 0, but in practice you could also use the mean
X_num = X_num.fillna(0)
```

```

# 3) Standardize data types for Keras.
# - Force all features to float32 (recommended format for neural networks)
X_num = X_num.astype('float32')

# 4) Convert the target variable (y) to numeric as well.
# - Any non-numeric values are coerced to NaN, then replaced with 0.
# - Finally, cast to int32 (since this is a classification target).
y_num = pd.to_numeric(y, errors='coerce').fillna(0).astype('int32')

X_train = X_train.apply(pd.to_numeric, errors='coerce').fillna(0).astype('float32')
X_test = X_test.apply(pd.to_numeric, errors='coerce').fillna(0).astype('float32')
y_train = pd.to_numeric(y_train, errors='coerce').fillna(0).astype('int32')
y_test = pd.to_numeric(y_test, errors='coerce').fillna(0).astype('int32')

print("Object dtype columns in X_train:",
      list(X_train.columns[X_train.dtypes == 'object']))
print("Any NaNs? ->", X_train.isna().any().any(), y_train.isna().any())

n_features = X_train.shape[1]
X_train_seq = np.asarray(X_train, dtype=np.float32).reshape(-1, n_features,
X_test_seq = np.asarray(X_test, dtype=np.float32).reshape(-1, n_features,

```

Object dtype columns in X_train: []

Any NaNs? -> False False

```

In [152]: model_lstm = Sequential([
    LSTM(32, input_shape=(X_train_seq.shape[1], X_train_seq.shape[2])),
    Dense(1, activation='sigmoid')
])
model_lstm.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])

# 6) Eğitim
hist_lstm = model_lstm.fit(
    X_train_seq, y_train,
    epochs=20, batch_size=32,
    validation_split=0.2, verbose=1
)

# 7) Tahmin ve metrikler
proba_lstm = model_lstm.predict(X_test_seq).ravel()
y_pred_lstm = (proba_lstm >= 0.5).astype(int)


# Accuracy ve AUC (bunlar zaten sende var)
print("Accuracy:", accuracy_score(y_test, y_pred_lstm))
print("AUC:", roc_auc_score(y_test, proba_lstm))


# Recall ve F1'i ek import olmadan hesapla (class=1)
cm = confusion_matrix(y_test, y_pred_lstm, labels=[0, 1])
tn, fp, fn, tp = cm.ravel()
precision = tp / (tp + fp) if (tp + fp) > 0 else 0.0
recall = tp / (tp + fn) if (tp + fn) > 0 else 0.0
f1 = (2 * precision * recall / (precision + recall)) if (precision + recall) > 0 else 0.0


print("Recall (class=1):", round(recall, 4))
print("F1 (class=1):", round(f1, 4))


```


```
# Detaylı rapor ve matris (bunlar da sende var)
print("\nClassification Report:\n", classification_report(y_test, y_pred_lst
print("Confusion Matrix:\n", cm)
```


Epoch 1/20
19/19  1s 20ms/step - accuracy: 0.5580 - loss: 0.6878 - val_accuracy: 0.5714 - val_loss: 0.6789


Epoch 2/20
19/19  0s 9ms/step - accuracy: 0.5580 - loss: 0.6695 - val_accuracy: 0.5782 - val_loss: 0.6606


Epoch 3/20
19/19  0s 9ms/step - accuracy: 0.5956 - loss: 0.6558 - val_accuracy: 0.6122 - val_loss: 0.6503

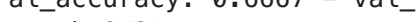
Epoch 4/20
19/19  0s 9ms/step - accuracy: 0.6246 - loss: 0.6434 - val_accuracy: 0.6599 - val_loss: 0.6427

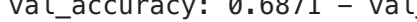
Epoch 5/20
19/19  0s 9ms/step - accuracy: 0.6314 - loss: 0.6317 - val_accuracy: 0.6395 - val_loss: 0.6400


Epoch 6/20
19/19  0s 9ms/step - accuracy: 0.6365 - loss: 0.6253 - val_accuracy: 0.6395 - val_loss: 0.6356

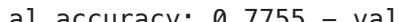
Epoch 7/20
19/19  0s 9ms/step - accuracy: 0.6451 - loss: 0.6176 - val_accuracy: 0.6667 - val_loss: 0.6280


Epoch 8/20
19/19  0s 9ms/step - accuracy: 0.6519 - loss: 0.6068 - val_accuracy: 0.6667 - val_loss: 0.6158


Epoch 9/20
19/19  0s 10ms/step - accuracy: 0.6621 - loss: 0.5901 - val_accuracy: 0.6871 - val_loss: 0.5978


Epoch 10/20
19/19  0s 9ms/step - accuracy: 0.7048 - loss: 0.5663 - val_accuracy: 0.7143 - val_loss: 0.5764


Epoch 11/20
19/19  0s 9ms/step - accuracy: 0.7611 - loss: 0.5396 - val_accuracy: 0.7755 - val_loss: 0.5569


Epoch 12/20
19/19  0s 9ms/step - accuracy: 0.8089 - loss: 0.5152 - val_accuracy: 0.7823 - val_loss: 0.5415


Epoch 13/20
19/19  0s 9ms/step - accuracy: 0.8259 - loss: 0.4964 - val_accuracy: 0.7891 - val_loss: 0.5313


Epoch 14/20
19/19  0s 9ms/step - accuracy: 0.8276 - loss: 0.4825 - val_accuracy: 0.7959 - val_loss: 0.5247

Epoch 15/20
19/19  0s 9ms/step - accuracy: 0.8294 - loss: 0.4726 - val_accuracy: 0.7959 - val_loss: 0.5209

Epoch 16/20
19/19  0s 9ms/step - accuracy: 0.8294 - loss: 0.4652 - val_accuracy: 0.7959 - val_loss: 0.5192

Epoch 17/20
19/19  0s 11ms/step - accuracy: 0.8276 - loss: 0.4597 - val_accuracy: 0.7959 - val_loss: 0.5187

Epoch 18/20
19/19  0s 9ms/step - accuracy: 0.8259 - loss: 0.4555 - val_accuracy: 0.7959 - val_loss: 0.5187

Epoch 19/20
19/19  0s 9ms/step - accuracy: 0.8276 - loss: 0.4517 - val_accuracy: 0.7959 - val_loss: 0.5187

al_accuracy: 0.7959 - val_loss: 0.5187

Epoch 20/20

19/19  0s 9ms/step - accuracy: 0.8294 - loss: 0.4480 - v

al_accuracy: 0.7959 - val_loss: 0.5187

6/6  0s 12ms/step

Accuracy: 0.7934782608695652

AUC: 0.8950263032042085

Recall (class=1): 0.8039

F1 (class=1): 0.8119

Classification Report:

	precision	recall	f1-score	support
0	0.76	0.78	0.77	82
1	0.82	0.80	0.81	102
accuracy			0.79	184
macro avg	0.79	0.79	0.79	184
weighted avg	0.79	0.79	0.79	184

Confusion Matrix:


[[64 18]

[20 82]]

```
In [156... # Deep Learning model performance results
data_dl = {
    "Model": ["Baseline MLP", "Enhanced MLP", "LSTM"],
    "Accuracy": [0.88, 0.87, 0.80],
    "Precision": [0.91, 0.88, 0.82],
    "Recall": [0.86, 0.88, 0.80],
    "F1-score": [0.88, 0.88, 0.81],
    "ROC AUC": [0.93, 0.94, 0.90]
}

# Convert to DataFrame
df_dl = pd.DataFrame(data_dl).set_index("Model")

# Plot heatmap (red tones)
plt.figure(figsize=(8, 5))
sns.heatmap(df_dl, annot=True, fmt=".2f", cmap="Reds", vmin=0.75, vmax=0.98,
plt.title("Deep Learning Model Performance Comparison", fontsize=14)
plt.yticks(rotation=0)
plt.show()
```

```
In [158... # Deep Learning model performance results
data_dl = {
    "Model": [
        "Baseline MLP",
        "Enhanced MLP",
        "LSTM"
    ],
    "Accuracy": [0.88, 0.87, 0.80],
    "Precision": [0.91, 0.88, 0.82],
    "Recall": [0.86, 0.88, 0.80],
    "F1-score": [0.88, 0.88, 0.81],
    "ROC AUC": [0.93, 0.94, 0.90]
}

df_dl = pd.DataFrame(data_dl)
print(df_dl)
```

	Model	Accuracy	Precision	Recall	F1-score	ROC AUC
0	Baseline MLP	0.88	0.91	0.86	0.88	0.93
1	Enhanced MLP	0.87	0.88	0.88	0.88	0.94
2	LSTM	0.80	0.82	0.80	0.81	0.90

Thank you